

Saad Zubairi

saadhzubairi@outlook.com | (646) 232-7806 | github.com/saadhzubairi | [linkedin.com/in/saadhzubairi](https://www.linkedin.com/in/saadhzubairi) | saadhzubairi.netlify.app

Summary

MS Computer Engineering at NYU (May 2026). Full stack engineer who ships production systems. Recent outcomes: cut ticket triage time 60-70% with a Zendesk extension (92% auto tag accuracy on 1.2k tickets/week), modernized NYU's legacy exam portal (T3 stack, 4 tier RBAC) serving 2.5k+ submissions/term, and built an MLflow tracked LSTM+TCN pipeline for crop yield forecasting. Targeting SWE/ML/Platform roles.

Education

New York University, United States — *MS, Computer Engineering* **Expected Graduation: May 2026**
Institute of Business Administration, Karachi, Pakistan — *BS, Computer Science* **Graduated: May 2023**

Relevant Coursework: Digital Signal Processing, Deep Learning, Internet Architecture, Distributed Systems, Object Oriented Programming, Data Structures, Algorithms, Operating Systems, Database Systems, Software Engineering

Experience

New York University, New York City — *Full Stack Engineer* **Jan 2025 – Present**

- Engineered a modernized exam portal (T3 stack) serving 2.5k+ submissions/term; reduced load times by 50% and improved user experience with a responsive design.
- Built scheduling/grading flows and admin dashboards, with email notifications and PDF reporting; eliminated double bookings (99% reduction) and improved auditability with event logs
- Implemented a comprehensive RBAC system with role-based access control, ensuring secure and efficient data handling.
- Enhanced platform reliability and scalability by implementing robust monitoring and alerting (Prometheus, Grafana)
- Automated CI/CD pipelines with GitHub Actions, and containerized deployments using Docker. Integrated statistical reporting and analytics dashboards for real-time insights into system usage and performance.

Cloaked, New York City — *Software Engineering Intern* **Jun 2025 – Aug 2025**

- Built a data pipeline that aggregated and geocoded mailbox locations from 5+ providers via headless scraping.
- Designed a daily status pipeline (<2% failures) and dashboard (Next.js/OpenLayers) with one-click CSV/PDF exports.
- Shipped a RAG-powered Zendesk auto-tagging Chrome extension that processed ~1,200 tickets/week with 92% tagging accuracy (P=0.92, R=0.88; n=1,238).
- Cut average triage time from 6 hours to 2.5 hours by automating categorization and tagging.
- Developed an MVP educational tool in Nuxt.js, delivering modular interactive learning components via a generative-UI.
- Built prompt-eval harnesses with precision/recall logging, preventing regressions.

Sadiq.ai, Santa Clara — *Software Engineer* **May 2024 – Aug 2024**

- Led a 5-engineer scrum team and shipped the SupportSphere MVP in 8 weeks, running weekly clinician feedback demos.
- Designed PostgreSQL schema with migrations + core data models and GraphQL based API endpoints in Nest.js.
- Built a robust RBAC system in Nest.js, supporting granular permissions and secure multi-role workflows.
- Integrated Redis for caching and session management, improving API response times and scalability under load.
- Improved reliability with Prometheus/Grafana, CI/CD, Docker; added real-time analytics.
- Provisioned infra (Terraform, 6 EC2/RDS) + automated deploys (GitHub Actions).

United Bank Limited, Karachi — *Software Engineer* **July 2023 – Apr 2024**

- Authored and migrated 12 account-management and 18 payment reconciliation endpoints in Temenos Fabric, reducing integration lead time by 60% and lowering API error rate to <0.5%.
- Designed and maintained all multi-threaded Java services within the Temenos Infinity platform.
- Developed and enhanced frontend elements for the bank's consumer based application using Visualizer, collaborating with UI/UX teams to improve user workflows and interface responsiveness.
- Implemented features; bug fixes across Kotlin UI and Java services.

Projects

SparseDeep EEG/ECG	Thesis: model-based deep learning for EEG/ECG analysis by integrating sparse signal priors (algorithm unrolling) into neural networks for interpretable, efficient, robust signal processing
AgriYieldPredictor	ML pipeline for multi-crop yield prediction: ingests HRRR/USDA/Sentinel data, trains LSTM+TCN models, and serves reproducible inference via MLflow + OpenStack Swift
TalentHive	MERN-stack web app with Python NER resume parser and TF-IDF/cosine-similarity job matching; matched 100+ profiles to jobs with 78% relevance in a 20-user pilot

Skills

Languages:	Python, TypeScript, Java, Go, SQL, C++, JavaScript, Rust, Kotlin
AI/ML:	PyTorch, LangChain, RAG, Vector Databases, LLM Fine-tuning, MLflow
Cloud & DevOps:	AWS, Docker, Terraform, Ansible, Linux, Git, GitHub Actions, OpenStack Swift
Frameworks:	Next.js, Nuxt.js, Express.js, Nest.js, Django, FastAPI, T3 Stack, REST, GraphQL, Node.js
Databases:	PostgreSQL, MongoDB, Redis, Vector Databases, MySQL
Frontend & Mobile:	React.js, Vue.js, OpenLayers, Android (Kotlin), Chrome Extensions
Monitoring & Tools:	Prometheus, Grafana, Selenium